

Model Inference Latency by Budget (averaged across datasets)

