


Accuracy vs Inference Latency
(all model×dataset×budget)



Accuracy vs Latency (averaged across datasets)
by model & budget

